

CRAL-TRIO: mean loss learning curve across seeds
descriptors_lipprop | class weights | lr=0.0001 | weight_decay=0.01 | final_m=4 | final_dropout=0.1 | hiddim=8 | heads=8 | batch=16 | seeds=0,1,2,3,4

