

IP\_trans: mean loss learning curve across seeds  
descriptors\_mlp\_coarse | class weights | lr=0.0001 | weight\_decay=0.01 | final\_m=4 | final\_dropout=0.1 | hiddim=8 | heads=8 | batch=16 | seeds=0,1,2,3,4

