

CRAL-TRIO: mean balanced_accuracy learning curve across seeds

descriptors_no_extent_coarse | class weights | lr=0.0001 | weight_decay=0.01 | final_m=4 | final_dropout=0.1 | hiddim=8 | heads=8 | batch=16 | seeds=0,1,2,3,4

