


IP_trans: mean loss learning curve across seeds

descriptors_no_extent_coarse_add_hid24 | class weights | lr=0.0001 | weight_decay=0.01 | final_m=4 | final_dropout=0.1 | hiddim=24 | heads=8 | batch=16 | seeds=0,1,2,3,4



best shared BA epoch 60
BA window mean=0.6278
seeds>=5
loss mean=0.6467