

IP_trans: mean precision learning curve across seeds

structural_pretrain_family_scratch | class weights | lr=0.0001 | weight_decay=0.001 | final_m=4 | final_dropout=0.1 | hiddim=64 | heads=8 | batch=16 | seeds=0,1,2,3,4

